

## 1 General Formula

**Frobenius Norm:**  $\|\theta\|_F = \sqrt{\sum_j \theta_j^2}$  **L1 Norm:**  $\sum_i |\theta_i|$

## 2 Neurons

### 2.1 Biology

Neuron contains **Soma**, **Axon**, **Dendrites**. Signals travel away from Soma via Axon.

#### 2.1.1 Membrane

Membranes contain Sodium and Potassium Channels and a Sodium-Potassium Pump. Each Channel:

1. Na: Outside of cell
2. K: Inside of cell
3. Na-K Pump: 3 Sodium out for 2 Potassium in

**Voltage Difference (Membrane Potential):** Difference in voltage on either side of membrane. Resting potential is  $-70\text{mV}$ . This resting potential is enforced by the Na-K Pump.

**Action Potential:** Electrical impulse travelling along axon to the synapse.

#### 2.2 Hodgkin-Huxley Model


Model of Neuron based on Ion Channel components.

Fraction of K+ channels open is  $n(t)^4$ . There are four identical gates in K+ channel. Fraction of Na+ channels open is  $h(t)m(t)^3$ . There are three identical gates and one unique gate in Na+ channel.

Gate dynamic system is a(t):  $\frac{da}{dt} = \frac{1}{\tau_a(V)}(a_{\infty}(V) - a)$

Membrane Potential Dynamics:  $\frac{dV}{dt} = J_{in} - g_L(V - V_L) - g_{Na}m^3h(V - V_{Na}) - g_Kn^4(V - V_K)$  c: Membrane capacitance  $I = C \frac{dV}{dt}$ : Net current inside cell  $J_{in}$ : Input current from other Neurons  $g_L$ : Leak conductance (membrane not impenetrable to ions)  $g_{Na}$ : Maximum Na conductance  $g_K$ : Maximum K conductance

Action Potential form:



Process: Stimulus breaks threshold causing Na+ channels to open, then close at action potential. Potassium channels open at action potential and close at refractory period.

### 2.3 Leaky-Integrate-and-Fire Model

Spike shape is constant over time, more important to know when spiked than shape. LIF only considers sub-threshold voltage and when peaked.

Dynamics system:  $c \frac{dV}{dt} = J_{in} - g_L(V - V_L) \Rightarrow \tau_m \frac{dV}{dt} = V_{in} - (V - V_L)$  for  $V < V_L$

This is dimensioned. Dimensionless converts  $v_{in} = \frac{v_{in} - V_L}{V_L - V_L}$  to become  $\tau_m \frac{dv}{dt} = v_{in} - v$

Then spike occurs when  $v = 1$  and we set a refractory period of  $t_{ref}$  before starting at 0 again.

Explicit Model:  $v(t) = v_{in}(1 - e^{-\frac{t}{\tau}})$

**Firing Rate:**  $\frac{1}{\tau_{ref} - \tau_m \ln(1 - v_c)}$  for  $v_c > 1$  **Tuning Curve:** Graph showing how neuron reacts to different input currents.

## 2.4 Activation Functions (Sigmoidal)

Logistic Curve:  $\frac{1}{1+e^{-x}}$  Arctan:  $\arctan(z)$  Hyperbolic Tangent:  $\tanh(z)$  Threshold: 0 if  $z < 0$ ; 1 if  $z \geq 0$  Rectified Linear Unit:  $\max(0, z)$  Softplus:  $\log(1 + e^z)$

## 2.4.1 Multi-Neuron Activation Functions

**SoftMax:**  $\frac{\exp(z_i)}{\sum_j \exp(z_j)}$  Converts elements to probability distribution (sum to 1) and normalizes values **ArgMax:** Largest element remains nonzero, everything else 0

### 3 Synapses

In real neuron, **Presynaptic action potential** releases **neurotransmitters** across **synaptic cleft** which binds to **postsynaptic receptors**

Equation for current entering **postsynaptic neuron:**  $h(t) = \begin{cases} kt^n e^{-\frac{t}{\tau}} & \text{if } t \geq 0 \text{ (for some } n \in \mathbb{Z}_{>0}) \\ 0 & \text{if } t < 0 \end{cases}$

$k$  is selected so  $\int_0^\infty h(t)dt = 1 \Rightarrow k = \frac{1}{n\tau^n + 1}$

**Postsynaptic potential filter:**  $h(t)$

**Spike train:** Series of multiple spikes  $a(t) = \sum_{p=1}^k \delta(t - t_p)$

**Dirac function:** Infinite at  $t = 0$ , 0 everywhere else. Properties:  $\int_{-\infty}^\infty \delta(t)dt = 1$ ,  $\int_{-\infty}^\infty f(t)\delta(t - \tau)dt = f(\tau)$

**Filtered Spike Train:**  $s(t) = a(t) * h(t)$  For each spike in spike train, run the postsynaptic potential filter on it, then sum each spike. Essentially, convolve the spike train to the postsynaptic potential filter.

Derivation:

$$s(t) = a(t) * h(t) = \int \sum_p \delta(t - t_p) h(t - \tau) d\tau = \sum h(t - t_p)$$

### 3.1 Neuron Activities

Neurons have multiple connections, with different strengths (weights). We can represent weights between neuron layers via weight matrix:  $W \in \mathbb{R}^{N \times M}$

Compute neuron function as:  $\vec{z} = \vec{x}W + \vec{b}$ , thus  $\vec{y} = \sigma(\vec{z})$

#### 3.1.1 Bias Representation

Add a neuron with constant value 1 for each layer, and use its weights as biases.

$$\vec{x}W + \vec{b} = (\vec{x} \cdot 1) \cdot \begin{pmatrix} W \\ \vec{b} \end{pmatrix}$$

#### 3.1.2 Connections between spiking Neurons

Let  $n_i = 0$  for simplicity for  $h(t)$ , then it is a solution of  $\tau_s \frac{da}{dt} = -s$ ,  $s_i(0) = \frac{1}{\tau_s}$

### 3.2 Full LIF Model

Dynamics are described by:  $\begin{cases} \tau_m \frac{dv}{dt} = s_i - v_i & \text{if not in refractory period} \\ \tau_s \frac{ds_i}{dt} = -s_i \end{cases}$

If  $v_i$  reaches 1, then start refractory period, send spike, reset  $v_i$  to 0. If spike arrives from neuron  $j$ , then  $s_i \leftarrow s_i + \frac{w_{ij}}{\tau_s}$

## 4 Neural Learning

### 4.1 Supervised Learning

Desired output is known, and we can minimize error within our model's predictions

**Regression:** Output is continuous-based function of inputs, goal is to get closest to output function.

**Classification:** Outputs are in discrete categories.

### 4.2 Unsupervised Learning

Output is not known, so goal is to find efficient structure of input

### 4.3 Reinforcement Learning

Feedback is given, but it's uninformative, and depends on lots of variables (ex. Chess AI)

## 4.4 Optimization

Given neural model, goal is to optimize weights to minimize loss function.

$$\min_{\theta} \mathbb{E}_{(x,t) \in \text{data}} [\mathcal{L}(f(x; \theta), t(x))]$$

## 5 Universal Approximation Theorem

For all continuous functions in the domain of  $n$  parameters in  $[0, 1]$  domain each, can be approximated as finite sums of sigmoid functions.

**Sigmoid Function:** Goes to 1 for positive infinity and 0 for negative infinity.

### 5.1 Process

1. By giving infinite weight to  $\sigma(wx)$ , this approximates a threshold function.
2. We create piece function as difference of threshold functions:  $P(x; b, \delta) = H(x; b) - H(x; b + \delta)$
3. Approximate each section of the function, as  $G(x) = \sum_{j=1}^{N'} f(a_j) P(x; b_j, \delta_j)$  each is within  $\varepsilon$  band

## 6 Loss functions

Single Error:  $L(y, t)$  is error between one output  $y$  and target  $t$   
Dataset Error:  $\mathbb{E} = \frac{1}{N} \sum_{i=1}^N L(y_i, t_i)$  as average error over entire dataset.

### 6.1 Mean Squared Error

Single Error:  $L(y, t) = \frac{1}{2} \|y - t\|_2^2$  Activation Function: Linear activation functions (ReLU) Problems: linear regression problems.

### 6.2 Bernoulli Cross Entropy

Single Error:  $L(y, t) = -\log(P(y, t)) = -\log(t^y(1-y)^{1-y}) = -(t \log(y) + (1-t) \log(1-y))$  Activation Function: Logistic function Problems: Output values are in range  $[0, 1]$ ;

### 6.3 Categorical Cross Entropy

Single Error:  $L(y, t) = -\log(\text{pin}C_k, t)) = -\log\left(\prod_{k=1}^K (y^k)^{t^k}\right) = -\sum_{k=1}^K t^k \log(y^k)$  Activation Function: Softmax Problems: Classification problems with one-hot vectors

## 7 Gradient Descent


Our goal is to minimize  $E(\theta)$  (expected Error), so we define gradient  $\nabla_{\theta} E = (\frac{\partial E}{\partial \theta_1} \dots \frac{\partial E}{\partial \theta_n})$

Error gradient:  $\tau \frac{d\theta}{dt} = \nabla E(\theta)$ ,  $t$  is a parameter for "time" as we move through parameter space. Euler step:  $\theta_{n+1} = \theta_n + \kappa \nabla E(\theta_n)$

## 8 Error Backpropagation

$\nabla_{z^{(l+1)}} E = \frac{\partial E}{\partial z^{(l+1)}} h^{(l+1)} = \sigma'(z^{(l+1)}) = \sigma'(W^{(l)} h^{(l)} + b^{(l+1)})$  Basically,  $h$  is hidden layer,  $z$  is input current,  $W$  is weight matrix,  $b$  is bias.  $\nabla_{z^{(l)}} E = \frac{dh^{(l)}}{dz^{(l)}} \odot \left[ \nabla_{z^{(l+1)}} E \cdot (W^{(l)})^T \right] \odot$  is hadamard product, which does element by element multiplication. We transpose because  $W_{ij}$  is connection from  $i$ th node in  $l$  to  $j$ th node in  $l+1$

Note that  $\vec{a} = \vec{x}W$  in this diagram:



$$\frac{\partial E}{\partial W_{ij}^{(l)}} = \frac{\partial E}{\partial z_j^{(l+1)}} \frac{\partial z_j^{(l+1)}}{\partial W_{ij}^{(l)}} = \frac{\partial E}{\partial z_j^{(l+1)}} h_i^{(l)}$$

Finally,  $\nabla_{z^{(l)}} E = \sigma'(z^{(l)}) \odot \left( \nabla_{z^{(l+1)}} E \cdot (W^{(l)})^T \right) \nabla_{W^{(l)}} E = [h^{(l)}]^T \nabla_{z^{(l+1)}} E$

### 8.1 Vectorization

We can generalize this process to take a batch of samples by letting  $x$  be a matrix of samples instead of just one sample. Then, note  $\nabla_{z^{(l)}} E$  is a matrix with same dimension as  $z^{(l)}$  as desired. Further, note that  $\nabla_{W^{(l)}} E$  is a gradient vector that sums the weight gradient matrix from each sample.

## 9 Auto-Differentiation

### 9.1 Expression Graph

Each operation is a square with its variable dependencies. Each variable has a pointer to its creator, which is the operation that created it.

#### 9.1.1 Pseudocode

Function  $f = g(x, y, \dots)$   
• Create  $g$  Op object  
• Save references to args  $x, y, \dots$   
• Create Var for output  $f$   
• Set  $g.val$  as  $g(x, y, \dots)$   
• Set  $f.creator$  to the  $g$  Op.

### 9.2 Differentiate

With each object in our graph, store derivative of total expression with respect to itself in member  $grad$  Use chain rule with parent operation Op.grad to get current grad. Ex: If  $y$  is parent of  $x$ , then  $x.grad = y.grad \cdot \frac{\partial y}{\partial x}$  Also add wherever multiple branches converge, as is normal in derivatives calculation.

Backward method:

```
class Var:
    def backward(s):
        self.grad += s
        self.creator.backward(s)
```

```
class Op:
    def backward(s):
        for x in self.args:
            x.backward(s * partialDeriv(Op, x))
```

In Var, self.val, self.grad, s must have same shape. In Op, s must be shape of operation output

## 10 Neural Nets w/w Auto-Diff

### 10.1 Gradient Descent Pseudocode

- Initialize  $v, \kappa$
- Make expression graph for  $E$
- Until convergence:
  - Evaluate  $E$  at  $v$
  - Zero-grad
  - Calculate gradients
  - Update  $v \leftarrow v - \kappa \cdot v.grad$

### 10.2 Neural Learning

Optimizing our weights and biases for our loss function.  $W \leftarrow -\kappa \nabla_W E$  By making network with AD classes, we leverage backward() to optimize gradient computation.

#### 10.2.1 Pseudocode

Given Dataset  $(X, T)$  and network model **net**, with parameters  $\theta$  and loss function  $L$

- $y = \text{net}(X)$
- $\text{loss} = L(y, T)$
- $\text{loss.zero\_grad}()$
- $\text{loss.backward}()$
- $\theta \leftarrow \kappa \cdot \theta.grad$

## 11 Overfitting

If big discrepancy in accuracy between training and testing, then we're "overfitting"

### 11.1 Problems

If training is too small, it's overfitting. If test error is much bigger than training error.

#### 11.1.1 Solutions

**Validation:** Keep subsection of testing error as validation error to determine proper hyperparameters, then finally we test to determine the correctness. **Regularization by Weight Decay:** Expand error to worry about hyperparameters as well.  $\tilde{E}(\hat{y}, t; \theta) = E(\hat{y}, t; \theta) + \frac{\lambda}{2} \|\theta\|_2^2$  Then new rules are  $\nabla_{\theta_i} \tilde{E} = \nabla_{\theta_i} E + \lambda \theta_i$  and  $\theta_i \leftarrow -\kappa \nabla_{\theta_i} \tilde{E} = -\kappa \lambda \theta_i$  **Data Augmentation:** Add slightly modified versions of data to make more samples **Dropout:** Drop some random nodes to distribute computation.

## 12 Optimization

### 12.1 Stochastic Gradient Descent

Take random batch of samples, then take the expected value over all of them:  $E(\hat{y}, \hat{\tau}) = \frac{1}{B} \sum_{d=1}^B (y_{\tau_d}, t_{\tau_d})$

### 12.2 Momentum

Usecases: When gradient descent oscillates very often before converging to a local minimum; optimization stops at shallow local optimum without global optimum. Gradient is a force instead of slope.  $\theta_{n+1} = \theta_n + \Delta t v_n$  and  $v_{n+1} = (1 - r) v_n + \Delta t A_n$  where  $A_n$  represents the gradient vector, so we make  $v^{(t)} \leftarrow \beta v^{(t-1)} + \nabla_W E$  and  $w^{(t)} \leftarrow w^{(t-1)} - \eta v^{(t)}$